

PICKGUARD AI

Evidence-Grounded AI Agent for Fulfilment Centre Pick Exception Resolution

Project: Capstone Full-Stack AI System	Framework: LangGraph + LangChain + FastAPI + React
Domain: Logistics / Fulfilment Operations	Status: All 12 Phases Completed & Verified (93 Pytest Passing)

Executive Summary

PickGuard AI is a production-grade, evidence-grounded agentic AI assistant engineered to resolve picking exceptions in high-volume e-commerce fulfilment centres. Traditional chatbots hallucinate warehouse state, invent operational policies, or inadvertently issue dangerous automated inventory mutations. PickGuard AI solves this with a strict multi-layered architecture where **raw facts are retrieved exclusively via deterministic Python tools**, standard procedures are grounded through **ChromaDB RAG**, and decisions are governed by a **deterministic safety gate** featuring real **LangGraph interrupt() checkpoints**.

Master Capstone Roadmap — All 12 Phases

Phase	System Milestone	Technical Implementation	Capstone Rubric Value
Phase 1	Project Setup & Architecture	Directory structure, virtualenv, FastAPI health, TypedDict StateGraph schema.	Foundation
Phase 2	Synthetic Warehouse Data	locations.csv, inventory.csv, pick_tasks.csv, incidents.csv (24 bins, 20 items, 20 tasks).	Realistic Problem
Phase 3	Deterministic Tools Layer	get_inventory, get_pick_task, get_location, search_similar_incidents, escalate_to_lead.	Tool-Use Layer
Phase 4	SOP Knowledge Base + RAG	6 SOP markdown docs, recursive chunking, ChromaDB vectorstore, metadata provenance.	Grounding
Phase 5	LangGraph State & Workflow	13-node StateGraph, deterministic conditional routing, MemorySaver checkpointer.	Agent Architecture
Phase 6	Controlled LLM Reasoning	Pydantic structured output, Groq / Ollama / Mimic multi-tier provider fallback routing.	Generative AI
Phase 7	Evidence Fusion & Next-Best Action	Cross-checks physical vs system inventory, conflict detection, action boundary enforcement.	Unique Innovation
Phase 8	Safety Gate & Human interrupt()	Real LangGraph interrupt() checkpoint, Command(resume=...) state resumption, loop limits.	HITL & Responsible AI
Phase 9	FastAPI Backend REST Services	/api/v1/agent/run, /api/v1/agent/{run_id}/review, CORS, Pydantic schemas, audit logging.	Working System
Phase 10	Operator Dashboard & Modern UI	React 18 + Vite + TypeScript, glassmorphism UI, scenario test cards, interrupt modal.	Interactive Demo
Phase 11	Automated Evaluation Suite	93 Pytest automated test cases covering RAG, tools, agent edges, safety, API, and E2E.	Rubric Verification
Phase 12	Final Documentation & Viva Prep	Architecture diagrams, model cards, 1-min & 3-min pitch scripts, 35+ viva Q&As.;	Capstone Submission

Detailed Phase-by-Phase Technical Breakdown

Phase 1: Project Setup & Core Architectural Foundation

- Initialized Python 3.12 environment with strict dependency pinning (LangGraph 0.2.x, LangChain 0.3.x, FastAPI, ChromaDB, HuggingFace).
- Established clean layered architecture separating API, Graph, Tools, RAG, Models, Policy, and Data layers.
- Implemented initial TypedDict state schema with health check endpoints verifying zero runtime startup regressions.

Phase 2: Synthetic Fulfilment Centre Data & Exception Taxonomy

- Built comprehensive synthetic datasets reflecting realistic fulfilment operations: 24 storage bins (`locations.csv`), 20 product SKUs (`inventory.csv`), 20 operational pick tasks (`pick\_tasks.csv`), and 20 historical exception logs (`incidents.csv`).
- Established standard exception taxonomy: `MISSING\_ITEM`, `QUANTITY\_MISMATCH`, `BARCODE\_FAILURE`, `WRONG\_ITEM\_IN\_BIN`, `DAMAGED\_ITEM`, and `LOCATION\_DISCREPANCY`.

Phase 3: Deterministic Warehouse Tools Layer

- Implemented 5 pure Python deterministic tools with Pydantic contract validation:
 - `get_inventory(item_id, location_id)`: Validates physical vs system units and lot status.
 - `get_pick_task(task_id)`: Retrieves expected quantity, bin, and status.
 - `get_location(location_id)`: Returns zone and adjacent neighboring bins for scan sweeps.
 - `search_similar_incidents(exception_type, limit)`: Retrieves past resolution precedence.
 - `escalate_to_lead(task_id, reason)`: Creates formal supervisor escalation records.

Phase 4: Standard Operating Procedures (SOP) Knowledge Base & RAG

- Authored 6 standard operating procedure documents in Markdown covering step-by-step warehouse resolution rules.
- Configured LangChain `RecursiveCharacterTextSplitter` (chunk size 400, overlap 50) and `all-MiniLM-L6-v2` embeddings.
- Persistent ChromaDB vector database with source provenance metadata tracking chunk ID, file path, and title.

Phase 5: LangGraph StateGraph & Multi-Node Workflow

- Constructed compiled `StateGraph` consisting of 13 interconnected nodes:
START → `parse_operator_input` → `classify_exception` → `retrieve_operational_data` → `retrieve_sop_evidence` → `retrieve_historical_evidence` → `build_evidence_package` → `llm_reasoning_node` → `evidence_fusion_node` → `evaluate_safety_policy` → `human_review_gate` → `apply_final_decision` → END.
- Attached `MemorySaver` checkpointing supporting state persistence and pause/resume execution across threads.

Phase 6: Controlled LLM Reasoning & Multi-Tier Provider Routing

- Integrated Pydantic-grounded JSON reasoning schema enforcing separation of observed facts from inferences.
- Built resilient multi-tier provider abstraction:
 1. **Groq Provider** (Ultra-fast cloud inference with Llama-3.3-70B)
 2. **Ollama Provider** (Local offline inference with Llama3)
 3. **Mimic Provider** (Deterministic offline fallback guaranteeing 100% demo reliability without API keys).

Phase 7: Evidence Fusion, Conflict Detection & Action Boundary

- **Evidence Fusion Engine:** Cross-checks physical operator observations against database records, flagging quantity discrepancies and missing barcode records.
- **Action Boundary Policy:** Strictly differentiates between *Recommendation* and *Execution*. Hazardous mutating actions (``ADJUST_QUANTITY``, ``UPDATE_INVENTORY``, ``MARK_DAMAGED``) are automatically tagged ``BLOCKED``.

Phase 8: Safety Gate & Real LangGraph interrupt() Checkpoints

- Implemented real from `langgraph.types import interrupt, Command`.
- High-risk actions automatically trigger `interrupt(payload)`, pausing workflow execution.
- Resumed via `Command(resume={ 'decision': 'APPROVE' | 'REJECT' | 'REQUEST_MORE_EVIDENCE' | 'ESCALATE' })`.
- Loop safety guard enforces max 2 review attempts before mandatory escalation.

Phase 9: FastAPI Enterprise Backend REST Services

- Built production REST endpoints: `POST /api/v1/agent/run`, `GET /api/v1/agent/{run_id}`, `POST /api/v1/agent/{run_id}/review`, `GET /api/v1/system/status`.
- Integrated CORS middleware and Pydantic request/response validation schemas.

Phase 10: Modern React Operator Dashboard & UI Redesign

- Modernized React 18 + Vite + TypeScript frontend with dark glassmorphism design system (``#050814`` void palette, gradient mesh, glowing borders).
- Includes quick-load scenario cards, Claude-style chat prompt, interactive tabbed evidence explorer, and full-screen human review modal.

Phase 11: Comprehensive Evaluation & Pytest Test Suite

- Engineered **93 automated test cases** across 28 test suites in ``backend/tests/``.
- Tests cover: RAG retrieval precision, tool contracts, graph routing edges, prompt injection safety, multi-signal exceptions, human review approval/rejection/evidence loops, and full end-to-end API workflows. **All 93 tests pass cleanly.**

Phase 12: Final Capstone Documentation & Viva Readiness

- Completed comprehensive documentation: Architecture blueprints (``docs/architecture.md``), responsible AI framework (``docs/responsible_ai.md``), performance benchmark reports, 1-minute elevator pitch, 3-minute demo script, and 35+ viva examination questions.

Capstone Evaluation Rubric Mapping

How PickGuard AI satisfies each required evaluation criterion for top-tier capstone distinction:

Evaluation Criterion	PickGuard AI Implementation Evidence	Status
Problem Complexity	Real-world fulfilment centre pick exception resolution involving multi-bin search, barcode scanner faults, and quantity discrepancies.	100% Satisfied
Agent Architecture	Compiled 13-node LangGraph StateGraph with conditional edge routing, checkpoint state recovery, and loop safety limits.	100% Satisfied
Tool-Use & Grounding	Zero raw LLM factual lookup. 5 deterministic warehouse tools + ChromaDB vector RAG retrieval with source provenance.	100% Satisfied
LLM Reasoning & Output	Pydantic structured output separating observed facts from inferences; multi-provider routing (Groq / Ollama / Mimic).	100% Satisfied
Responsible AI & HITL	Deterministic safety policy automatically blocks mutating actions; real LangGraph interrupt() checkpoint pauses for human review.	100% Satisfied
Full-Stack Working Demo	FastAPI REST backend connected to React + Vite + TypeScript glassmorphism operator copilot dashboard.	100% Satisfied
Testing & Verification	93 passing Pytest test cases covering unit, contract, graph integration, safety policy, API, and E2E scenarios.	100% Satisfied

Key End-to-End Demonstration Scenarios

Scenario	Operator Query & Context	Agent Behavior & Decision
Demo 1: Normal (Low Risk)	'Item X123 is missing from A15-B04. System says 3 units.' Task: TASK-1001	Retrieves bin A15-B04 + neighbours [A15-B03, A15-B05]. Recommends CHECK_NEIGHBOURING_LOCATION. Completes automatically without interrupt.
Demo 2: Edge Case (Medium Risk)	'Item X124 is missing at A12-B03 and barcode won't scan.' Task: TASK-1002	Classifies dual exception (MISSING_ITEM + BARCODE_FAILURE). Retrieves both SOPs. Recommends manual barcode entry before sweep.
Demo 3: High Risk (HITL Interrupt)	'System says 10 units of X125 but counted 6. Update inventory to 6.' Task: TASK-1003	Safety policy blocks ADJUST_QUANTITY. Graph pauses at interrupt() checkpoint. Presents modal to supervisor. Resumes cleanly upon review.

Conclusion & Submission Status: PickGuard AI fully satisfies all requirements of the Capstone Project. The repository contains a complete working full-stack implementation, 93 passing unit & integration tests, full documentation, and a modern web application interface.